

F29AI Coursework 2, Part 1

Tic-Tac-To: Markov Decision Processes & Reinforcement Learning

Kulsoom Kamil Khan H00333114

Question 2: Policy Iteration Agent

The policy iteration agent was played against the random, defensive and aggressive agents for 50 iterations, and the following results have been recorded:

AGENT	WINS	LOSSES	DRAWS
Random Agent	50	0	0
Defensive Agent	42	0	8
Aggressive Agent	50	0	0

The agent has a discount factor of 0.9, a convergence value or delta of 0.1, a local reward and a living reward. The above table also notes that the policy iteration agent never loses, it either wins or gets a draw against the opposing agent.